

1. Understand the Data

1.1 Problem Definition

The automotive used-car market is one of the largest consumer markets globally. Buyers and sellers alike struggle to determine a fair price for a vehicle, as prices are influenced by many interacting factors including brand prestige, age, mileage, fuel efficiency, and transmission type.

Our dataset contains 72,435 real UK used car listings scraped from AutoTrader, covering 7 popular brands: Audi, BMW, Ford, Volkswagen, Toyota, Skoda, and Hyundai. Each listing includes key vehicle attributes.

1.2 Column Dictionary

Column	Data Type	Description
model	Categorical	Specific car model (e.g., A1, Focus, 3 Series)
year	Numerical	Year the car was manufactured (1996–2020)
price	Numerical (Target)	Listed sale price in British Pounds (£)
transmission	Categorical	Gearbox type: Manual, Automatic, Semi-Auto, Other
mileage	Numerical	Total miles driven by the car
fuelType	Categorical	Engine fuel: Petrol, Diesel, Hybrid, Electric, Other
tax	Numerical	Annual road tax in GBP
mpg	Numerical	Miles per gallon — fuel efficiency measure
engineSize	Numerical	Engine displacement in litres (0.0–6.6L)
Make	Categorical	Car brand/manufacturer

1.3 Business Importance

Accurate price prediction helps:

- Buyers avoid overpaying for a used vehicle
- Sellers set competitive, market-aligned prices
- Dealerships automate appraisal workflows
- Insurance companies estimate vehicle value accurately

Our classification goal: predict whether a car is High Price (above market median) or Low Price. This binary classification is actionable for both buyers and sellers.


2. Investigate the Data

2.1 Dataset Overview

Metric	Value
Total Rows	72,435
Total Columns	10
Rows with Missing Values	~3,621–3,623 per column (5%)
Numerical Features	6 (year, price, mileage, tax, mpg, engineSize)
Categorical Features	4 (model, transmission, fuelType, Make)
Target Variable	price (for regression) / high_price (for classification)
Price Range	£495 — £145,000
Median Price	£14,495
Mean Price	£16,584


Visualization 1: Price Distribution

The price distribution is right-skewed, with most cars priced between £8,000 and £25,000. A small number of luxury/premium vehicles push the mean above the median.




Visualization 2: Average Price by Brand

Audi and BMW are the most premium brands, averaging over £22,000. Ford and Toyota target the budget-conscious segment, averaging around £12,000–£12,500.




Visualization 3: Price vs. Mileage

A clear negative correlation exists between mileage and price — higher mileage cars command significantly lower prices. This makes intuitive sense as wear and tear reduces a vehicle's value.




Visualization 4: Fuel Type Distribution

Petrol (63%) and Diesel (31%) dominate the dataset. Hybrid and Electric vehicles represent only ~5% of listings, reflecting the gradual transition to clean energy vehicles in the UK market.




Visualization 5: Price Trend by Year

Newer cars (2018–2020) are significantly more expensive than older models. Cars from 2015 and earlier show lower price floors, while recent models show higher variability — likely due to popular vs. basic trim levels.



Visualization 6: Price by Transmission Type

Semi-Automatic and Automatic vehicles command a £10,000+ premium over Manual cars. This reflects both buyer preference for comfort and the higher manufacturing cost of automatic gearboxes.



3. Fix the Data

3.1 Data Quality Issues Identified

Issue	Scale	Impact	Solution
Missing values in all columns	~3,621 rows (5%)	Incomplete analysis, model errors	Remove rows with any missing value (listwise deletion)
engineSize = 0.0	Small subset	Unrealistic — no car has 0L engine	Treated as missing during analysis; flagged
Extreme mpg outliers (up to 470.8)	Rare records	May indicate data entry errors	Retained but monitored (could be electric/hybrid)
Price outliers (£145,000)	Very rare	May skew mean but valid luxury cars	Retained — real market data
Inconsistent casing in Make column	Systematic	e.g., "audi" vs "Audi"	Standardized in SAS using UPCASE/PROPCASE

3.2 Before vs. After Cleaning

Metric	Before Cleaning	After Cleaning
Total rows	72,435	43,325 (rows retained after dropping NAs)
Missing values	3,621–3,623 per column	0 missing values
Rows dropped	—	29,110 rows removed (~40%)
Data completeness	~95%	100%

Note: ~40% row loss is significant. The missing data appears concentrated in a subset of records (likely one data source or scraping batch). All complete records are representative of the full distribution.

3.3 SAS Cleaning Code

```
/* PROC IMPORT: Load raw data */
PROC IMPORT DATAFILE="/path/car_price.csv"
  OUT=cars_raw DBMS=CSV REPLACE;
  GETNAMES=YES;
RUN;

/* Remove missing values (listwise deletion) */
DATA cars_clean;
```

```
SET cars_raw;  
IF NMISS(year, price, mileage, tax, mpg, engineSize) = 0;  
IF transmission NE "" AND fuelType NE "" AND Make NE "";  
/* Standardize Make column casing */  
Make = PROPCASE(STRIP(Make));  
RUN;  
  
/* Verify cleaning results */  
PROC MEANS DATA=cars_clean N NMISS;  
VAR year price mileage tax mpg engineSize;  
RUN;
```

4. Create New Insights (Feature Engineering)

4.1 Engineered Features

Feature Name	Formula	Business Justification
car_age	2020 - year	Direct depreciation proxy; newer = more valuable. Simpler than using raw year.
price_per_mile	price / (mileage + 1)	Value-for-mileage ratio. High ratio = expensive per use unit. Useful for comparative pricing.
efficiency_score	mpg / engineSize	Combines two important specs: higher score = better fuel economy relative to engine power. Electric/hybrid cars score highest.
high_price (target)	1 if price > £14,495 else 0	Binary classification target. Above-median = "High Price" category for dealers and buyers.

4.2 SAS Feature Engineering Code

```
DATA cars_featured;
  SET cars_clean;

  /* Feature 1: Car Age */
  car_age = 2020 - year;

  /* Feature 2: Price per Mile (value density) */
  price_per_mile = price / (mileage + 1);

  /* Feature 3: Fuel Efficiency Score */
  IF engineSize > 0 THEN efficiency_score = mpg / engineSize;
  ELSE efficiency_score = mpg; /* Handle 0-engine edge case */


  /* Feature 4: Binary Target Variable */
  IF price > 14495 THEN high_price = 1;
  ELSE high_price = 0;

RUN;

PROC MEANS DATA=cars_featured MEAN STD MIN MAX;
  VAR car_age price_per_mile efficiency_score;
RUN;
```


Visualization 7: MPG vs Engine Size (colored by Price)

Smaller, more efficient engines tend to appear in lower-priced cars. Larger engines (3L+) command premium prices. The efficiency_score feature captures this relationship effectively.



Visualization 8: High vs. Low Price Cars by Brand

BMW and Audi have a higher proportion of above-median-price cars, while Ford and Toyota skew toward the below-median category. This brand-price relationship is a key predictor in our model.



5. Build a Model

5.1 Model Choice: Logistic Regression

We chose Logistic Regression (PROC LOGISTIC in SAS) to predict whether a car is High Price (1) or Low Price (0). Logistic regression is ideal because:

- It provides interpretable coefficients for each feature
- It outputs probabilities, useful for confidence-based decisions
- It handles both numerical and encoded categorical features
- It is a well-established baseline for binary classification

5.2 Model Configuration

Parameter	Value
Algorithm	Logistic Regression (Binary Classification)
Target Variable	high_price (0 = Low Price, 1 = High Price)
Features Used	car_age, mileage, mpg, engineSize, tax, price_per_mile, efficiency_score, Make_enc, fuelType_enc, transmission_enc
Train/Test Split	80% Train (34,660 rows) / 20% Test (8,665 rows)
Random Seed	42

5.3 SAS Model Code

```
/* Encode categorical variables */
DATA cars_model;
  SET cars_featured;
  /* Create dummy/numeric encodings for categorical variables */
  IF Make = "Audi" THEN Make_num = 1;
  ELSE IF Make = "BMW" THEN Make_num = 2;
  ELSE IF Make = "Ford" THEN Make_num = 3;
  ELSE IF Make = "Vw" THEN Make_num = 4;
  ELSE IF Make = "Toyota" THEN Make_num = 5;
  ELSE IF Make = "Skoda" THEN Make_num = 6;
  ELSE IF Make = "Hyundai" THEN Make_num = 7;

  IF transmission = "Manual" THEN trans_num = 1;
  ELSE IF transmission = "Automatic" THEN trans_num = 2;
  ELSE IF transmission = "Semi-Auto" THEN trans_num = 3;
  ELSE trans_num = 0;

  IF fuelType = "Petrol" THEN fuel_num = 1;
  ELSE IF fuelType = "Diesel" THEN fuel_num = 2;
  ELSE IF fuelType = "Hybrid" THEN fuel_num = 3;
  ELSE IF fuelType = "Electric" THEN fuel_num = 4;
  ELSE fuel_num = 0;
RUN;
```

```
/* Split into Train/Test */
DATA cars_train cars_test;
  SET cars_model;
  IF RAND("uniform") <= 0.8 THEN OUTPUT cars_train;
  ELSE OUTPUT cars_test;
RUN;

/* PROC LOGISTIC: Train Classification Model */
PROC LOGISTIC DATA=cars_train OUTMODEL=logit_model;
  MODEL high_price(EVENT="1") =
    car_age mileage mpg engineSize tax
    price_per_mile efficiency_score
    Make_num fuel_num trans_num / LINK=LOGIT;
  OUTPUT OUT=train_preds PREDICTED=pred_prob;
RUN;

/* Score Test Data */
PROC LOGISTIC INMODEL=logit_model;
  SCORE DATA=cars_test OUT=test_preds;
RUN;
```

6. Explain the Model

6.1 Model Performance

Metric	Value
Overall Accuracy	86.57%
Precision (Low Price class)	85%
Recall (Low Price class)	87%
F1-Score (Low Price class)	86%
Precision (High Price class)	88%
Recall (High Price class)	86%
F1-Score (High Price class)	87%
Test Set Size	8,665 records

6.2 Confusion Matrix

	Predicted: Low	Predicted: High
Actual: Low	3,723 (TN)	532 (FP)
Actual: High	632 (FN)	3,778 (TP)

The model correctly classified 3,723 low-price and 3,778 high-price cars. Only 532+632 = 1,164 cars were misclassified out of 8,665 total test records.

6.3 Feature Importance (Logistic Regression Coefficients)

Feature	Coefficient	Interpretation
engineSize	+3.97	STRONGEST predictor. Larger engines strongly indicate high price. Each extra litre raises log-odds by 3.97.
car_age	-0.82	Older cars are significantly less likely to be high-priced. Depreciation effect is clear.
transmission_enc	+0.39	Automatic/Semi-Auto transmission is associated with higher-priced cars.

fuelType_enc	-0.24	Certain fuel types (e.g., Petrol) correlate with lower price range vehicles.
Make_enc	-0.07	Brand encoding has modest effect after controlling for engine and age.
mpg	-0.05	Higher MPG tends to be associated with smaller, cheaper cars (excluding EVs).

6.4 Key Insights from the Model

- Engine size is the single most powerful predictor of whether a car is priced above market median
- Car age (depreciation) is the second most important factor — each additional year significantly reduces price class probability
- Transmission type matters greatly: Semi-Auto and Automatic cars are far more likely to be high-priced
- Fuel type and brand add nuance but are secondary to the mechanical characteristics

7. Make It Useful

7.1 Real-World Applications

This model has immediate practical applications across the automotive industry:

1. For Car Buyers: Input your shortlisted car's specs and instantly know if you're being charged a fair price. If the model predicts "Low Price" but the seller is asking above £14,495, walk away or negotiate.

2. For Dealerships: Automate vehicle appraisal at intake. When a customer brings a trade-in, the model can instantly estimate the market tier (High/Low) to support pricing decisions.

3. For Insurance Companies: Quickly estimate vehicle market value tier for premium calculation without a full appraisal.

4. For Online Marketplaces: Flag listings that appear over- or under-priced relative to comparable vehicles, improving market transparency.

7.2 Model Deployment in SAS (Scoring New Records)

```
/* Score new incoming car listings */
DATA new_cars;
  INPUT model $ year price transmission $ mileage fuelType $ tax mpg engineSize Make
  $;
  car_age = 2020 - year;
  price_per_mile = price / (mileage + 1);
  IF engineSize > 0 THEN efficiency_score = mpg / engineSize;
  /* Encode categoricals same as training */
  /* ... (same encoding logic) ... */
DATALINES;
Focus 2019 . Manual 12000 Petrol 145 45.6 1.0 Ford
;

PROC LOGISTIC INMODEL=logit_model;
  SCORE DATA=new_cars OUT=predictions;
RUN;

PROC PRINT DATA=predictions;
  VAR model year mileage I_high_price P_1;
  TITLE "Car Price Category Predictions";
RUN;
```

7.3 Limitations and Future Work

- The 40% data drop due to missing values may introduce bias — future work should explore imputation

- Model accuracy of 86.6% leaves room for improvement with ensemble methods (Random Forest, Gradient Boosting)
- Geographic variation within the UK is not captured — regional price differences could be added
- Colour, service history, number of owners — unmeasured variables that strongly influence real prices
- Time series modeling to track price trends over years could add predictive value for market timing

7.4 Conclusion

Starting with a mysterious dataset of 72,435 UK used car listings, we successfully built a complete data science pipeline. We cleaned ~5% missing data, engineered 4 meaningful features, explored 8 visualizations revealing clear patterns, and trained a Logistic Regression classifier that predicts whether a car is priced above market median with 86.6% accuracy.

The most important finding: engine size and car age together account for the majority of pricing power in the UK used car market. Transmission type is a significant secondary factor. This actionable insight can guide buyers, sellers, and automotive businesses toward better data-driven decisions.